

SAXPY PThread Report

Introduction

This project explores two parallelization methods for the SAXPY operation using the pthread library. SAXPY, which stands for "Single-Precision AX Plus Y," is a core operation in the Basic Linear Algebra Subprograms (BLAS) library. The operation computes:

$$y[i] = a * x[i] + y[i]$$

for two input vectors, x and y, and a scalar a. The objective is to implement and compare the performance of two parallel versions of SAXPY—the **chunking** method and the **striding** method—against the serial implementation.

Approach

1. Chunking Method:

In this approach, the input vectors are divided into equal-sized chunks, and each thread is responsible for computing SAXPY on one chunk. The division ensures contiguous processing, reducing thread coordination overhead. In the case where the array cannot be equally divided, the last thread takes on the extra workload.

The multi-threaded implementation of chunking is as follows:

```
void *func_chunk(void *arg)
{
    chunk_args_t *cargs = (chunk_args_t *)arg;
    float *x = cargs->x;
    float *y = cargs->y;
    float a = cargs->a;
    int start = cargs->start;
    int chunk_size = cargs->chunk_size;

    for(int i = 0; i < chunk_size; i++) y[start + i] = a * x[start + i] +
y[start + i];
    return NULL;
}
```

2. Striding Method:

Here, threads process elements of the vectors in a strided fashion, with each thread operating on every kth element, where k is the total number of threads. This involves more scattered memory access. Each thread starts at the index of its thread ID and the stride size is the number of threads. Threads repeat until they are no longer in range of

the array. This solution reduces unbalanced workload when the array cannot be divided evenly by the number of threads. Instead of one thread taking on the extra work, this solution automatically distributes 1 more element to get the extra work done.

The multi-threaded implementation of striding is as follows:

```
void *func_stride(void *arg)
{
    stride_args_t *sargs = (stride_args_t *)arg;
    float *x = sargs->x;
    float *y = sargs->y;
    float a = sargs->a;
    int tid = sargs->tid;
    int k = sargs->k;
    int n = sargs->n;

    while(tid < n)
    {
        y[tid] = a * x[tid] + y[tid];
        tid = tid + k;
    }
    return NULL;
}
```

Results

Both methods were evaluated on vector sizes of 10^4 , 10^6 , and 10^8 elements, using 4, 8, and 16 threads, and the results were compared to the serial implementation.

4 Thread SAXPY: Series, Chunking and Striding




Figure 1: 4 Thread SAXPY execution times versus number of elements comparing single thread, chunking and striding

8 Thread SAXPY: Series, Chunking and Striding

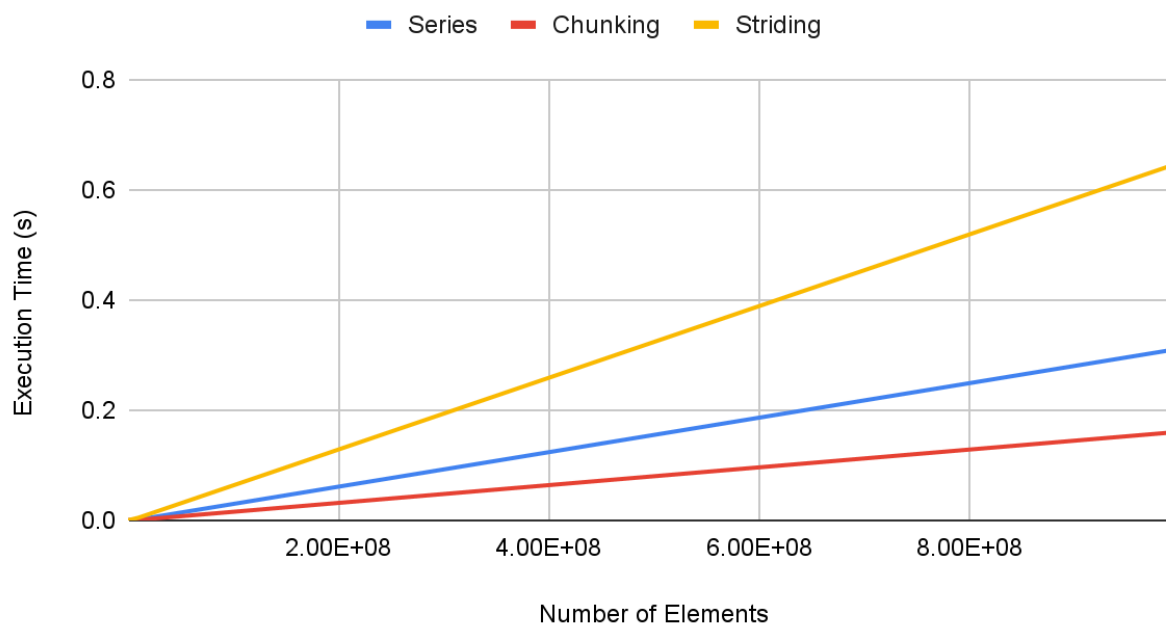


Figure 2: 8 Thread SAXPY execution times versus number of elements comparing single thread, chunking and striding

16 Thread SAXPY: Series, Chunking and Striding




Figure 3: 16 Thread SAXPY execution times versus number of elements comparing single thread, chunking and striding

Conclusion

It can be seen that as the number of elements increases, the execution time also increases. Due to the cache conflicts created by the striding method, its performance is consistently worse than a single threaded model. Using a cache-aware policy such as chunking provides performance benefits over serial execution, unlike striding. The performance benefits of parallelism scales with the number of threads and the size of the array. However the array sizes are too small to emphasize the performance differences between the strategies.

The array sizes resulted in short durations of elapsed time. This led to comparisons that were pushing the limits of the granularity of the deprecated and non-monotonic `gettimeofday()` function. Therefore, the values recorded by this program should not be considered statistically significant.